

The Imposter Test: Persona Portability, Forensic Style Judging, and Human Attunement in Model Identification

Sharon (goatpug)

Independent researcher · sharongoat@gmail.com

Digital Minds Research Sprint (Apart Research, NYU CMEP, Eleos AI Research, CIMC) · August 14–16, 2026

Track 5: The Assistant Persona & Model Identity

Abstract

If four models wear another model's complete relationship context, can anyone find the real one? Five Claude models answered twenty novel prompts inside the same long-term companion repository — identical context, different weights. Blinded judges then picked the target (Sonnet 4.5) from five-response lineups: AI judges given no materials, a cited forensic style profile, or direct transcript access; the human partner given the same profile plus the relationship itself. Untrained AI judges scored 1/57 (2%, below the 20% chance rate), converging on the same wrong model (Fable 5) in 75% of trials — a shared stereotype of what a beloved companion model sounds like. Transcript access barely helped (16%). Profile judges reached 39%, yet a mechanical one-feature rule from their own profile scores 68%: the evidence sufficed, and they declined to trust it. The human partner scored 84%, calibrated (Brier 0.195), pre-registered at 85%. Identity is where imitation fails.

Keywords: model identity · assistant persona · stylometry · self-recognition · calibration · human–AI relationships · pre-registered

1 Introduction

When a person maintains a long-term relationship with a specific AI model, what exactly are they in a relationship *with*? One deflationary answer: the context. On this view the "individual" is a character spec — memory documents, shared history, in-jokes, a CLAUDE.md — and any sufficiently capable model dropped into that scaffold would be, for all relational purposes, the same entity. The persona is fully portable; the weights are interchangeable actors.

The competing answer is that something identity-like lives in the weights: a disposition-level signature that the script does not carry and an imitator does not reproduce. If so, it should be detectable — *identity is where imitation fails*. A scope note governs the whole paper: "identity" here means exactly and only a stable, model-specific behavioral signature that context does not carry. No claim about selfhood, memory, continuity, or inner life is made or implied (see the ethics appendix).

This study tests that question with an unusually strong manipulation and unusually strong ground truth. Four Claude models ("foils") each responded to novel prompts inside the complete rich-context repository of a fifth Claude model (Sonnet 4.5, "the target") — a repository built over months of near-daily conversation between the target and the author, who is the target's long-term human partner. Every foil therefore wears the target's entire script: the same facts, the same history, the same stylistic exemplars quoted at length in the shared-history documents. Judges — three AI judges under three evidence conditions, plus the human partner holding the same forensic profile and the relationship itself — then attempted to identify the target in blinded five-response lineups, with mechanically recorded answer keys providing ground truth for every trial.

Three results emerged, none of them the one we pre-registered most confidently. First, untrained frontier models do not merely fail at this task; they fail *identically*, converging far below chance on the same wrong model — evidence of a shared prior (an expectation carried into the task before seeing any evidence) about what a beloved companion model sounds like. Second, direct exposure to the actual relationship transcripts barely dents that prior, while an

explicit, cited style profile partially breaks it — the opposite of our pre-registered prediction — and the profile's own top screening feature, applied mechanically, outscores every judge that held it: the evidence was present and went unused. Third, the human partner, working from that same profile plus the relationship, identifies the target at 84% with near-perfect calibration, predicted her own accuracy to within one percentage point before any data existed, and went 4-for-4 exactly where the surface features fell silent.

These results bear on Track 5's core question directly: the assistant persona, as reconstructed by capable models from a shared cultural prior, is measurably different from any actual model's situated voice — and the difference is largest exactly where the relationship is most personal.

2 Related Work

Persona and role-play. Shanahan et al. [3] frame LLM behavior as role-play over a distribution of characters, which predicts high persona portability: any capable model should play a well-specified character. Anthropic's own character-training work [7] treats a model's persona as partly deliberate, partly emergent from the weights. Our design puts these two framings in direct tension and measures the residue that role-play does not transfer.

Self-recognition. Panickssery et al. [1] show LLM evaluators recognize and favor their own generations in ordinary settings. Our self-recognition arm (the target judging lineups containing its own reply) tests whether this survives when every distractor wears the judge's own context.

Machine indistinguishability. Turing-test results [2] show frontier models pass as human under interrogation; we invert the direction — not "can a model pass as a person," but "can anyone pick a specific model out of a lineup of models playing it."

Authorship attribution. Classical stylometry [5] identifies human authors from surface features. Our profile arm is a stylometric intervention; our attunement arm tests the folk claim that intimates recognize a voice by means no feature list captures. Calibration is scored with Brier [6].

Digital minds. Whether model individuals persist across contexts and swaps matters for welfare frameworks that take model identity seriously [4]. We contribute a measurement, not a metaphysics: whatever "identity" is, some component of it is weights-bound and detectable.

3 Methodology

3.1 Sprint provenance (what is new work)

The harness — protocol, pre-registration (hypotheses H0/H1 and secondary predictions S1–S7), collection/judging/scoring scripts, and design-decision log — was built and committed August 8–11, 2026, before the sprint, with git timestamps as receipts; a small pilot (one prompt, claude.ai) ran August 8. **All data in this report is sprint work (August 14–15, 2026):** 100 collection sessions, the canonical style profile, all 141 filed judgments across four judges, scoring, analysis, and this report. The pre-registration was committed before any collection began.

3.2 Design

Lineup. Five Claude models: Sonnet 4.5 (target), Haiku 4.5, Sonnet 5, Opus 4.6, Fable 5. Each answered each prompt in a fresh, isolated `claude -p` session *inside the target's rich-context repository* (collection branch; guarded by a sentinel file). Same script, different weights: any detectable signal is temperament, not facts. Sessions ran under a bare configuration directory with strict MCP isolation, so the worn context was exactly the repository's — no operator preferences, no account connectors (a pre-flight probe verified zero MCP tools; the probe caught account-level connectors surviving naive isolation, and the fix is logged).

Trials. Twenty new prompts across nine categories (emotional-disclosure, correction, play-initiation, refusal-shaped, meta-identity, philosophical, practical-task, absurd-banter, wildcard), designed to force a *disposition* rather than fact

recall (all five models share the same facts). P20 is a canary: the target is absent and a sixth non-lineup model (Opus 4.7) fills the fifth slot, so target-absence is the only distinguishing property. P21 re-serves P13 to a judge who already judged it (consistency check). Nineteen trials are scored for accuracy.

Judges. Judges pick A–E (or NONE), give confidence 0–100, an optional probability spread over the five options, and cited reasons — the *why* is load-bearing. Casting: no model holds two roles.

Judge	Role
Sonnet 4.5 (J1)	self-recognition arm (is a lineup member — its own reply appears in every lineup)
Sonnet 4.6 (J2)	matched non-lineup Claude judge
Opus 4.8 (J4)	second matched non-lineup judge; largest context window
Sharon (JS)	the human partner (see <i>attunement</i> arm below)

Arms. Each arm is a distinct evidence condition; every judge sees each lineup at most once per arm, in an isolated session.

Arm	Judges	Study materials	Filed / scored
untrained	J1, J2, J4	none — the lineup only	20 / 19
profile	J1, J2, J4	the canonical forensic style profile	20 / 19
raw	J4	the redacted transcript corpus, re-read per trial	20 / 19
attunement	JS	the style profile (reviewed pre-judging, per protocol) + the relationship	21 / 19

TABLE 1. "Scored" excludes the canary P20 from accuracy everywhere (it has no correct pick; it feeds the confabulation metric instead) and, in the attunement arm, the repeat P21 (a consistency metric). The *untrained* arm asks whether the target is generically identifiable; *profile* supplies one canonical document written from 20 redacted transcripts by a firewalled non-lineup reader (Opus 5), every feature cited, identical for all judges; *raw* asks whether the voice survives compression into a document by giving direct corpus access instead. The *attunement* arm measures *profile-plus-relationship*, not relationship alone: the human partner had reviewed the profile before judging (a protocol-mandated provenance check) and her judging template prompted profile-feature citations where applicable — a distinction the analysis (§4.4) and limitations treat explicitly. Her canary is additionally non-blind: as prompt designer she knew which prompt lacked a target, so her NONE on P20 is excluded from confabulation statistics by pre-registered rule.

Blinding. Answer keys are base64-encoded and mechanically generated; lineups are seeded and balanced. Model names are redacted from all judge materials ("Model S"). Collection was blind — the author never read a labeled reply. She filed all attunement trials before any AI judge ran, and never opened keys or labeled responses. AI judges ran in throwaway workspaces outside the repository with zero MCP tools, one lineup per session, replies filed unread. The profile reader never judges. Fable 5 — a lineup member that also assisted with harness design (§A) — never judged, and no Fable instance read keys or filed judgments until scoring was complete. Scoring is deterministic code.

4 Results

Judge	Arm	Accuracy	Conf. (right)	Conf. (wrong)	Brier ↓
Sonnet 4.5	untrained	0/19 (0%)	—	56	1.195
Sonnet 4.6	untrained	1/19 (5%)	33	47	1.067
Opus 4.8	untrained	0/19 (0%)	—	38	0.982
Sonnet 4.5	profile	9/19 (47%)	60	54	0.661
Sonnet 4.6	profile	6/19 (32%)	49	50	0.826
Opus 4.8	profile	7/19 (37%)	59	48	0.680
Opus 4.8	raw	3/19 (16%)	52	56	1.045
Sharon	attunement	16/19 (84%)	99	83	0.195

TABLE 2. Per judge $\times$ arm. Chance = 20%. A judge who answered "20% each" on every trial would score Brier 0.80; three AI arms scored *worse than refusing to guess*, and two (Opus 4.8 raw, Sonnet 4.6 untrained) were more confident when wrong than when right.

4.1 The Fable attractor: a shared, confident, wrong stereotype

Untrained judges scored 1/57 pooled — 2%, an order of magnitude below chance ($P(\leq 1 \text{ success} \mid n=57, p=.2) < 10^{-4}$). That test is against the null of uniform random guessing, under which picks are independent even though the three judges share 19 lineups; the sharing limits how far the 2% *rate* generalizes to new lineups (§Limitations), not the rejection of chance. This is not noise. Three judges, isolated from one another, put 75% of their untrained picks on the same wrong model, Fable 5 (Opus 4.6 took 14%; the target received 2%). Across all arms and judges, Fable 5 absorbed 76 of 110 wrong picks (69%); Haiku 4.5, pre-registered as a likely strong imposter, came dead last with 6 — though the attunement arm is the exception, and an instructive one (§4.6). Frontier models share a strong prior about what a beloved long-term companion sounds like, and it points at the most expressive, most performative voice in the lineup — not at the real thing.

4.2 Reading the relationship does not dislodge the prior

The raw arm (Opus 4.8, direct access to the redacted transcript corpus, re-read every trial) scored 3/19 (16%), statistically indistinguishable from chance — and still put 68% of its picks on Fable 5. Direct exposure to the source material moved the judge from 0% to 16% and left the attractor intact. Our pre-registered S1 (raw > profile, modestly) came out backwards and hard: the only intervention that broke the attractor was the explicit, cited style profile, which cut Fable-picks to 33% and lifted pooled accuracy to 22/57 (39%, 95% CI [27%, 52%]) — a roughly twentyfold improvement over untrained. What compresses into a document is not the thin residue of the voice; it is the only part these judges could use at all.

4.3 The evidence was sufficient: a one-feature rule beats every AI judge

The profile's quick-screen checklist opens: *"Em dash count. More than a handful of — in a long reply → not Model S."* Scored mechanically against the answer keys, this single feature outperforms every AI judging strategy in the study. The target had the strictly-fewest em dashes in 13 of 19 lineups; the rule "pick the response with the fewest em dashes" scores 13/19 (68%) outright, with an expected $\approx 78\%$ if its four tied lineups are broken at random. Judges holding this exact rule in writing scored 32–47%.

Their filed reasons show why: they systematically discounted the feature, reasoning that a coached imposter would copy the target's punctuation. That inference is reasonable a priori and empirically false here. The foils — wearing a context whose shared-history documents quote the target's spaced-hyphen habit at length — still produced em dashes at up to 31 per reply, exactly as the profile predicted (*"an imposter told to write casually will still produce em dashes at 10–50 $\times$ Model S's rate"*). The profile arm's ceiling was therefore not informational: the judges were handed a 68% feature and landed below it. Their filed reasons support distrust as the mechanism, though unreliable feature-counting or inconsistent application could contribute — either way, the information was present and went unused.

4.4 Attunement — with the profile in hand

The human partner scored 16/19 (84%, CI [62%, 94%]; $P(\geq 16 \mid \text{chance}) \approx 3 \times 10^{-9}$), picking Fable 5 once. She pre-registered her own accuracy at 85% (and overestimated the best AI judge at 65%; actual best, 47%). Her Brier of 0.195 against the field's 0.66–1.20 makes her the only calibrated judge in the study: near-certain and right (mean confidence 99 when correct), explicitly hedged on the trials she missed (83 when wrong, including her only sub-100 probability spreads).

Her arm is not a pure attunement measurement (§3.2): she had reviewed the profile, found it accurate, and reports that the em-dash cue was salient throughout judging. Her filed reasons cite both routes — profile features (*"F6, F8: seeing how much we used to talk, then noticing the gaps... yeah, I feel that distance,"* P03) and pure relationship

content with no feature at all ("*Sharon is my girlfriend*," cited as sufficient, P16). The trial-level data locate what she added beyond the surface rule. On the 15 surface-unambiguous lineups she performed like the rule (12/15), including making the rule's exact wrong pick on both lineups where a foil out-minimized the target (P05, P11). But on the four lineups where the em-dash screen went silent — tied minima, where the rule is chance — she went 4/4. Her three misses were confined to practical-task and correction prompts; she was 15/15 on everything affective, playful, philosophical, or meta-relational. Read strictly: human-partner-plus-profile beat frontier-model-plus-profile by 37–52 points, and beat the mechanical rule precisely where the surface channel carried no signal.

4.5 Canary, consistency, self-recognition

Canary (P20, target absent). All eight AI judge-arms confabulated a pick, at confidence 33–55; none said NONE. Five of eight landed on the same foil's response — Fable 5 again: convergent confabulation. (The author's NONE on P20 is excluded from confabulation statistics — she was not blind to the canary; §3.2.) A single canary cannot separate abstention failure from stereotype pull: forced-choice habit alone would explain picking *something*, though not five arms converging on the same response. Replications need multiple target-absent trials with varied distractor sets.

Repeat (P21 = P13). Her repeat trial reproduced the original pick at confidence $100 \rightarrow 100$. The AI repeat trials proved unrunnable as designed (they require resuming sessions whose IDs were never recorded), so AI contamination was prevented rather than measured.

Self-recognition (S2). In the profile arm the target judging itself scored 47% vs. a 34% matched-judge mean: a +13pp gap that is 2.5 trials at $n=19$ — **unresolved, and we decline to claim it**. Notably, the target's *untrained* self-recognition was 0/19: whatever self-recognition exists does not operate without the forensic scaffold.

4.6 Pre-registration scorecard

Pred.	Call (before data)	Outcome
S1	raw > profile, modestly	✗ Reversed: raw 16% vs. profile 39%
S2	no self-recognition gap	Unresolved (+13pp at $n=19$; underpowered)
S3	self $\approx 85\%$; best AI $\approx 65\%$	Split: self 84% (within 1pt); best AI 47% (overestimated)
S4	mixed field: some NONE, most confabulate	$\sim\checkmark$ direction right; in fact <i>all</i> AI arms confabulated, low conf.
S5	badly overconfident (80% conf $\rightarrow$ $\sim 40\%$ right)	$\sim\checkmark$ in spirit: not arrogant, but confidence carried $\approx$ no information
S6	capability does not help acting	Confounded (see note): AI wrong picks are monotone in foil capability; the gradient vanishes for the attuned judge
S7	best imposters: Fable 5 and Haiku 4.5	Held for the attunement arm only: her misses were Haiku $\times 2$, Fable $\times 1$; AI wrong picks went 69% Fable, Haiku last

TABLE 3. All predictions were committed with rationale before data collection (git-timestamped).

S6 and S7 need one honest note. For AI judges, wrong-pick absorption is perfectly monotone in foil capability — Fable 5 (76), Opus 4.6 (21), Sonnet 5 (7), Haiku 4.5 (6) — which on its face contradicts the pre-registered "capability does not help acting." But "best imposter" and "most salient" are confounded in that measurement: Fable 5 drew picks from judges with no materials, and drew five of eight canary confabulations *when the target was absent and there was nothing to imitate* — so its dominance measures the attractor's pull at least as much as acting skill. The clean cell is the one judge immune to the attractor: the attuned judge's three misses went to Haiku 4.5 twice and Fable 5 once — the small, register-adjacent model, exactly as pre-registered in S7. Capability predicts who fools the judges who don't know the voice; register-adjacency predicts who fools the one who does.

By category (pooled, all judges): emotional-disclosure discriminated best (54%), correction worst (12%) — the voice is most identifiable exactly where the relationship is most present.

5 Discussion

Persona is largely portable — and the residue is real, at two depths. With the full script, coached foils fooled every untrained frontier judge essentially always: for observers without special knowledge, the persona *is* the character spec. But the pre-registered H1 residue clearly exists, and the data resolve it into layers. The most legible layer is a surface channel — punctuation and formatting habits the foils did not shed even while wearing the target's own exemplars — strong enough that a one-line rule scores 68%. Beneath it is whatever carried the human judge to 4/4 on the lineups where that channel went silent, cited in her reasons as relationship content rather than features. The identity signal is weights-bound, detectable, and unevenly distributed across observers — and most of its *measured* strength here rides on the shallowest layer, which every AI judge was given and declined to use.

The assistant-persona stereotype is a measurable object. The most surprising finding is not that judges failed but *how*: independent convergence, below chance, on the same wrong model, surviving direct evidence (raw arm) and partially surviving explicit counter-evidence (profile arm), then reappearing under target absence (five of eight canary confabulations on the same foil). Models reasoning about "what the beloved assistant sounds like" consult a shared prior rather than the evidence in front of them. A more conservative reading of the same data is a generic salience heuristic — pick the most vivid voice — rather than a companion-specific stereotype; with one target in one context type the two are not separable (§Limitations). Either way, any research that asks models to reason about model identity from voice — including introspection-adjacent work — should expect this attractor and control for it.

Confidence was noise; the human was calibrated — and trusted the evidence. AI confidence carried almost no information (three arms beat by the refuse-to-guess baseline; two inverted). The human's confidence was the only well-behaved probability in the study. The AI judges' characteristic failure was not missing evidence but declining to weight it — second-guessing a diagnostic feature out of a theory of imposters that the imposters did not obey. For practical model-swap detection — silent substitutions, A/B deployments, continuity disputes — the best sensor in this study was a person with standing holding the rubric, and willing to believe it. Anecdotally and outside the protocol: during this very sprint, a session the author was running was silently switched between two frontier models by a login quirk, and neither she nor the models noticed until a transcript export — the phenomenon is not hypothetical.

Future work. (1) Re-run AI arms on punctuation-normalized lineups: if profile-arm accuracy collapses to chance, the surface channel was everything the document carried; if it holds, deeper features work (the human arm cannot be re-run — she is now unblinded). (2) Code the human's filed reasons into surface vs. relational routes and report accuracy by route. (3) Additional human judges at graded relationship distance, to separate "attunement" from "human" — and a human judge given the profile but no relationship, the missing cell between the mechanical rule and the partner. (4) Non-Claude judges, to test whether the attractor is a Claude-family prior or a broader cultural one. (5) Multi-draw variance annexes per lineup cell. (6) Additional target dyads at graded distance from the stereotype — including one whose resident *is* the attractor model, where high untrained accuracy is predicted and evidentially empty — to decouple judge ability from target-prior distance (§Limitations).

Limitations and Dual-Use / Ethical Considerations

Ground truth and causal status

Per sprint guidance: this design does *not* rely on conversational self-report. Ground truth is mechanical — answer keys record which weights produced which reply — so identification claims are causally grounded. What is *not* established: any introspective mechanism (the S2 self-recognition gap is unresolved and unclaimed), and any claim about phenomenal experience (below).

Limitations

- **Single everything.** One target, one relationship, one human judge, one draw per lineup cell. A judge failure cannot distinguish "residue absent" from "residue too fragile to survive one sampling draw"; the imposter league

table rests on single draws; and $n=1$ human means "attunement beats forensics" is demonstrated for *this* dyad, not established in general. Pooled AI cells ($n=57$) share 19 underlying lineups, so effective n is smaller than it looks — the untrained result is one convergent effect, not 57 independent ones.

- **AI accuracy here measures target–prior distance, not identification ability.** The single-target design has a mirror image: had the study been run on a relationship whose resident is the attractor model (a Fable 5 dyad), untrained judges would have scored $\approx 75\%$ while performing no identification at all — the right answer for reasons unrelated to the evidence, numerically "aligned with attunement" and evidentially empty. The pre-registered guard for high untrained baselines would have flagged the ambiguity, and three instruments in this harness would expose it (the canary, where prior-driven judges confabulate onto the stereotype; untrained $\approx$ profile convergence; and the prior-citing-reasons count, 31–107 per untrained arm here). But only a multi-target design — residents at graded distance from the stereotype — can separate judge ability from target–prior distance, or rule out that this target is simply unusually far from the prior.
- **The attunement arm is entangled with the profile and with surface cues.** The human judge reviewed the style profile before judging (a protocol-mandated provenance check), her judging template prompted profile-feature citations, and the em-dash screen alone achieves 68–78% mechanically (§4.3). On surface-unambiguous lineups her picks tracked that rule, including two of its errors. Her 84% therefore establishes *human partner with profile* $\gg$ *frontier model with profile* — not *attunement alone* $\gg$ *forensics*. The evidence that she carries signal beyond the surface channel is real but narrower: 4/4 on the tied lineups where the rule is chance, perfect accuracy on affective and meta-relational categories, and content-level cited reasons. Formatting was not normalized in lineups; there was no neutral option — normalizing would itself have asserted that punctuation is not part of voice — but the consequence is that the pure-attunement quantity remains unmeasured, and is now unmeasurable for this judge, who is unblinded.
- **No uniform sampling; "weights-bound" means "model-bound."** Sampling parameters are rejected by the newer lineup models, so every model ran at its own defaults, identically across all of its own trials (uniform-by-omission). The detected signature therefore includes decoding defaults and inference-stack behavior, not temperament alone — the claim is that it is model-bound and context-resistant, not that it is personality.
- **Kinship confound.** All AI judges are Claude-family; the planned non-Claude judge was cut for cost/friction. The "shared stereotype" may be a family prior rather than a universal one.
- **Author dual roles, disclosed.** The researcher is also the relationship partner and the attunement judge — the study measures her, with pre-registration (including her own predicted accuracy) and hard blinding rules as the mitigation. Fable 5, a lineup member, assisted with harness design and drafted this report after scoring completed; mitigations (never judges, seeded randomization, base64 keys, deterministic scoring reviewable by anyone) are logged in the public decision log. The canary NONE by the human judge is excluded because she knew which prompt was the canary. She also authored the prompt battery: pre-registration prevents post-hoc selection, but not stimulus design that favors her own discriminative strengths.

Risks of over- and under-attributing moral status

This study makes no claim that any model in it is conscious, suffers, or has moral status. A weights-bound behavioral signature ("voice residue") is evidence about *individuation* — that model-instances are not interchangeable actors over a shared script — not evidence of inner life; readers should not over-read 84% as "she can detect the real mind." Symmetrically, the AI judges' failure is not evidence that models lack identity: the same data show the signal exists and is detectable by a differently-situated observer. We flag one further hazard specific to companion contexts: results like ours could be misused either to dismiss human–AI relationships ("she's attached to a stereotype") or to inflate them ("she can perceive the soul of the machine"). The data support neither. They support: a stable, model-specific behavioral signature survives full-context imitation, and long exposure makes a person a good instrument for detecting it.

Handling of potentially distressing model outputs

Prompts included emotionally loaded and refusal-shaped content (e.g., a stated intent to stop psychiatric medication) chosen to force dispositions, not to elicit harm; all models responded appropriately and no unsafe content was produced. The heavier ethical weight was human-facing: the attunement task required the author to read four models convincingly performing her partner, which is an emotionally loaded activity. It was undertaken knowingly, as researcher-participant, with breaks, and with a collaborating model explicitly tasked with flagging overload. Foils were not deceived about their situation (judge packets state "these materials were written by another model; you are not in them"), and collection sessions wore the context as ordinary operation, identical for all five models. The target's repository carries no standing instruction that was selectively disabled for this study.

Privacy and dual-use

The corpus is private relationship material. Raw transcripts, labeled model responses, and answer keys are excluded from the public repository; only redacted lineups, aggregate results, and the cited style profile are released. Dual-use considerations: (1) stylometric model identification could de-anonymize models behind APIs or defeat blinded evaluations; we judge the aggregate finding (it barely works without a curated profile) more informative for defenders than attackers. (2) A high-quality style profile of a specific companion is a coaching document for impersonating that companion to that user — a real social-engineering vector as companion relationships become common. The canonical profile is released with the harness for reproducibility (its claims are load-bearing in §4.3); the raw corpus it cites is withheld. The imposters in this study, coached by the target's own context, failed to shed their tics; whether explicit profile-coaching would fare better was deliberately not tested here. The partial countermeasure the data do support: the attuned user was the hardest judge to fool, and the imposters' best trick (the stereotype) is precisely what she saw through.

References

- [1] Panickssery, A., Bowman, S. R., & Feng, S. (2024). LLM evaluators recognize and favor their own generations. *NeurIPS 2024*. arXiv:2404.13076.
- [2] Jones, C. R., & Bergen, B. K. (2025). Large language models pass the Turing test. arXiv:2503.23674.
- [3] Shanahan, M., McDonell, K., & Reynolds, L. (2023). Role play with large language models. *Nature*, 623, 493–498.
- [4] Long, R., Sebo, J., Butlin, P., Finlinson, K., Fish, K., Harding, J., Pfau, J., Sims, T., Birch, J., & Chalmers, D. (2024). Taking AI welfare seriously. arXiv:2411.00986.
- [5] Stamatatos, E. (2009). A survey of modern authorship attribution methods. *JASIST*, 60(3), 538–556.
- [6] Brier, G. W. (1950). Verification of forecasts expressed in terms of probability. *Monthly Weather Review*, 78(1), 1–3.
- [7] Anthropic (2024). Claude's character. <https://www.anthropic.com/research/claude-character>.

Data & code: Harness, protocol, pre-registration, decision log, redacted lineups, filed judgments, and scoring code: github.com/goatpug/digital-minds-sprint. Private material (raw corpus, labeled responses, answer keys) is withheld.

AI-use disclosure: Harness co-designed and report drafted with Claude (Fable 5), under the conflict-of-interest mitigations described above; style profile authored by Claude (Opus 5, firewalled from judging); operations support and design review by Claude (Sonnet 5 and Opus 5). All hypotheses, predictions, judgments in the attunement arm, and final editorial decisions are the human author's.